



ElasticSearch

May 2017

Course Map - Searching and Analyzing Data with Elasticsearch

- 1 Overview
- 2 History of Search
- 3 How Does Search Works?
- 4 Inverted Index
- 5 Introducing Elasticsearch
- 6 Index, Shards, Replicas

Overview

- A little search engine history and the importance of search
- Basics steps involved in indexing and searching documents
- The inverted index, the heart of a search engine
- An introduction to Elasticsearch and its basic building blocks
- Set up and install Elasticsearch on your local machine and check cluster health

Overview

Prerequisites

- Familiarity with the command line on a Mac, Linux or Windows machine
- Familiarity with using RESTful APIs to perform actions
- A very basic understanding of distributed computing

Install and Setup

- The latest version of Elasticsearch, 5.4.0 requires Java version 8
- A Mac, Linux or Windows machine on which Elasticsearch can be installed

Course Overview

- **Introduction** to basic concepts in Elasticsearch, download and install
- **Building** an index, **adding** documents to it both individually and in bulk
- **Search** queries on an index using the Query DSL
- **Analysis** of data on an index using aggregations

Brief History of Search

1945

Vannevar Bush first talks of the need to index records

1991

Tim Berners-Lee combined hypertext, TCP and DNS to imagine WWW

1993

Excite improved search by using statistical analysis of word relationships

1970s

The ARPANet network which laid the foundation of the modern internet

1993

Primitive search engines, linear search of URLs, very basic ranking

1994

Yahoo offered a directory of useful webpages i.e. a portal

Brief History of Search

1994

Lycos provided ranking relevance, prefix matching, a huge catalog

1996

Inktomi pioneered the paid inclusion model

1998

Google ranking pages based on how many other pages link to it

1994

Altavista had natural language queries, inbound link checking

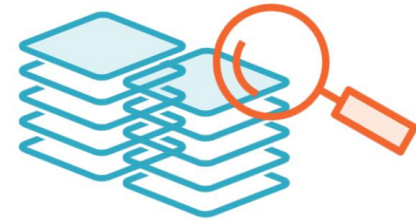
1997

ask.com had natural language search, human editors for queries

Today

Google, Bing, Baidu, Naver, Yahoo

How Does Search Work?



- **What Is the Objective of Search?**
 - Find the most relevant documents with your search terms

Most Relevant Document for Search Terms:



**Know of the
document's
existence**



**Index the
document for
lookup**



**Know how
relevant the
document is**



**Retrieve
ranked by
relevance**

How Does Search Work?

Most Relevant Document for Search Terms



Web crawler



**Inverted
index**



Scoring



Search

How Does Search Work?

Search is not restricted to the Web
Sites Have Their Own Search



E-commerce



Video



E-learning

Inverted Index

Documents Have Content

House Stark

Winter is coming

House Baratheon

Ours is the fury

House Tyrell

Growing Strong

Inverted Index

Tokenize Text into Words

winter
is
coming
ours
the
fury
growing
strong

split words

lowercased

removed
punctuation

Inverted Index

Tokenize Text into Words

winter	1
is	2
coming	1
ours	1
the	1
fury	1
growing	1
strong	1

Inverted Index

Tokenize Text into Words

winter	1
is	2
coming	1
ours	1
the	1
fury	1
growing	1
strong	1

Inverted Index

Tokenize Text into Words

winter	1	Stark
is	2	Stark, Baratheon
coming	1	Stark
ours	1	Baratheon
the	1	Baratheon
fury	1	Baratheon
growing	1	Tyrell
strong	1	Tyrell

Inverted Index

Tokenize Text into Words

winter	1	Stark
is	2	Stark, Baratheon
coming	1	Stark
ours	1	Baratheon
the	1	Baratheon
fury	1	Baratheon
growing	1	Tyrell
strong	1	Tyrell

Inverted Index

Dictionary

coming	1
fury	1
growing	1
is	2
ours	1
strong	1
the	1
winter	1

sorted so
lookup is easy

Stark
Baratheon
Tyrell
Stark, Baratheon
Baratheon
Tyrell
Baratheon
Stark

Inverted Index

Postings

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

Inverted Index

Search

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

winter

Inverted Index

Search

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

fury

Inverted Index

Search

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

is

Search

Inverted Index

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

coming OR strong

Search

Inverted Index

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

fury and growing

Search

Inverted Index

Searches Using Inverted Indices

- Find all words ending with “ong”

strong → gnorts

- Search for all words starting with “gno”

Search

Inverted Index

Searches Using Inverted Indices

- Split words into n-grams for substring search
- yours → yo, you, our, ours, urs
- Match substrings with n-grams

Searches Using Inverted Indices

- Geo-hashes for geographical search
- Algorithms such as Metaphone for phonetic matching
- “Did you mean?” searches use a Levenshtein automaton

Inverted Index

- An inverted index is at the heart of a search engine

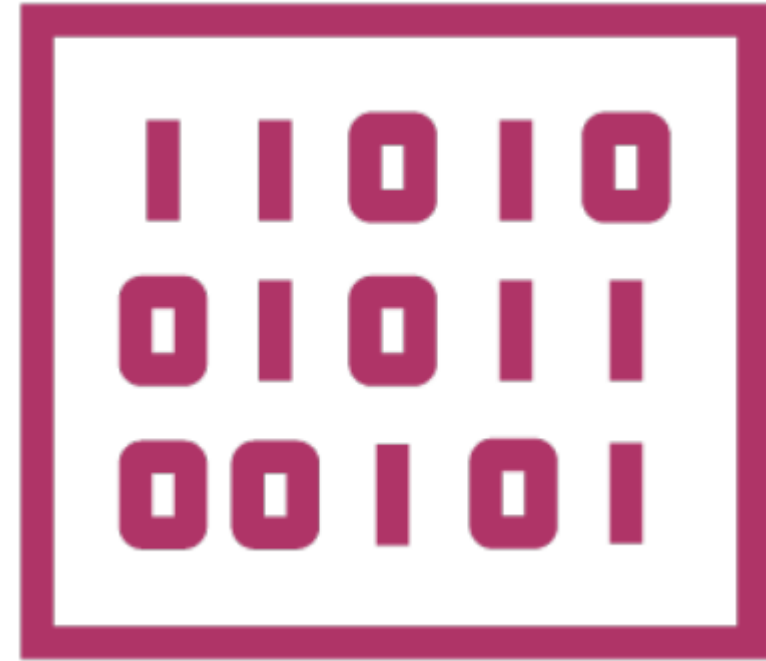
Implementing Search - Apache Lucene

Apache Lucene

The indexing and search library for a high performance, full-text search engine.

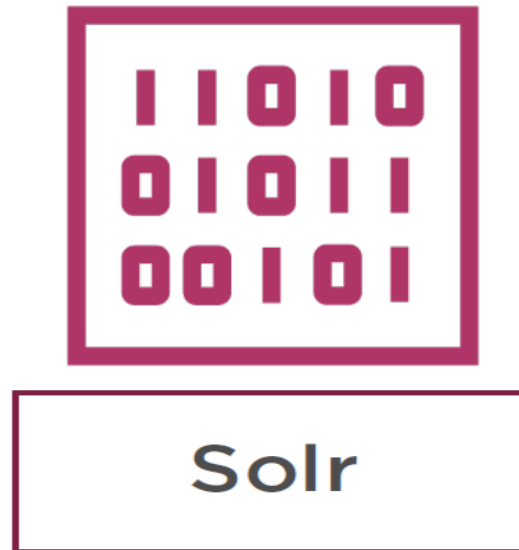
Open source, free to use written in Java, ported to other languages.

Just like Hadoop in the distributed computing world, Lucene is the nucleus of several technologies built around it.



Implementing Search - Apache Lucene

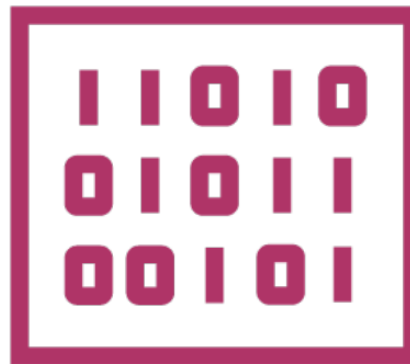
Apache Lucene



Web crawling and index parsing

Implementing Search - Apache Lucene

Apache Lucene

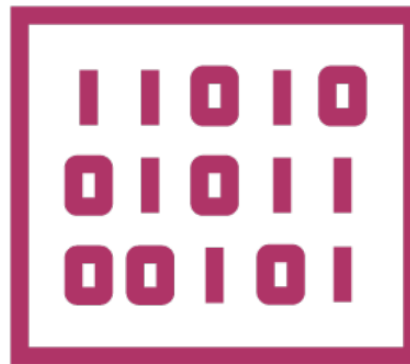


Nutch

Open source, free to use written in Java, ported to other languages

Implementing Search - Apache Lucene

Apache Lucene



Open source, SQL distributed database

Elasticsearch



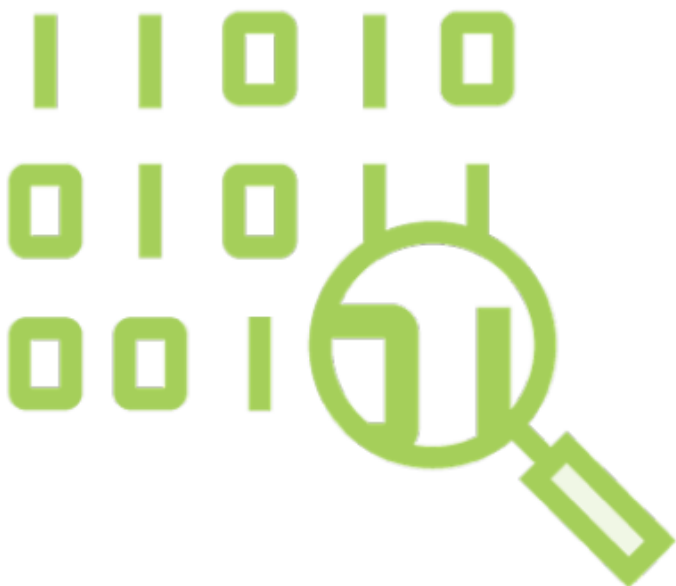
Elasticsearch is a distributed search and analytics engine which runs on Lucene

Introducing Elasticsearch



- An open source, search and analytics engine, written in Java built on Apache Lucene

Introducing Elasticsearch



- **Distributed**: Scales to thousands of nodes
- **High availability**: Multiple copies of data
- **RESTful API**: CRUD, monitoring and other operation via simple JSON-based HTTP calls
- **Powerful Query DSL**: Express complex queries simply
- **Schemaless**: Index data without an explicit schema

Elasticsearch



Product catalog
Inventory
Autocomplete



Video clips
Categories
Tags



Courses
Authors
Topics

Elasticsearch



**Mining log data
for insights**



**Price alerting
platform**



**Business analytics
and intelligence**

Working with Elasticsearch



**As a service in the
cloud**

<https://www.elastic.co/cloud/as-a-service>



**On your local
machine**

Basic Concepts of Elasticsearch

Near Realtime Search

Very low latency, **~1 second** from the time a document is **indexed** until it becomes **searchable**



Basic Concepts of Elasticsearch

Node

Single server
Performs **indexing**
Allows **search**
Has a **unique id**
and name



Basic Concepts of Elasticsearch

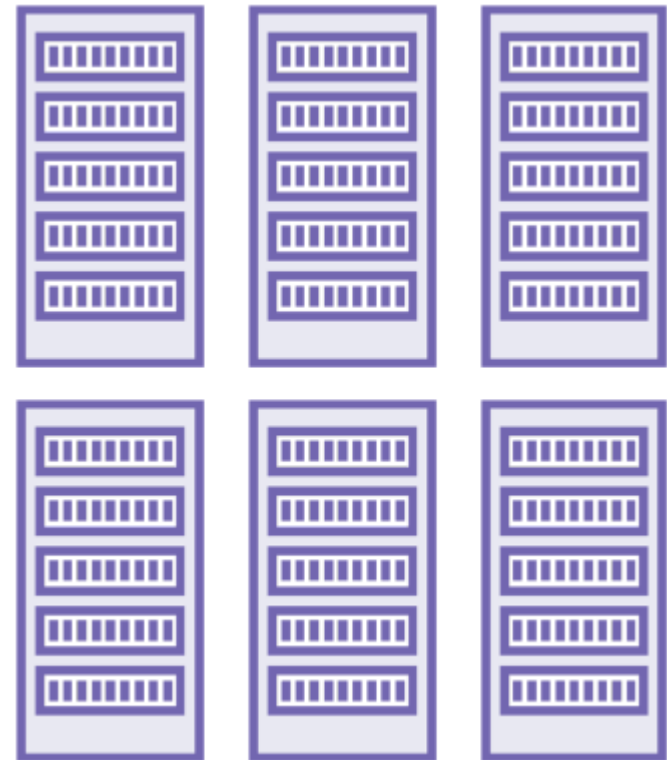
Cluster

Collection of nodes

Holds the entire
indexed data

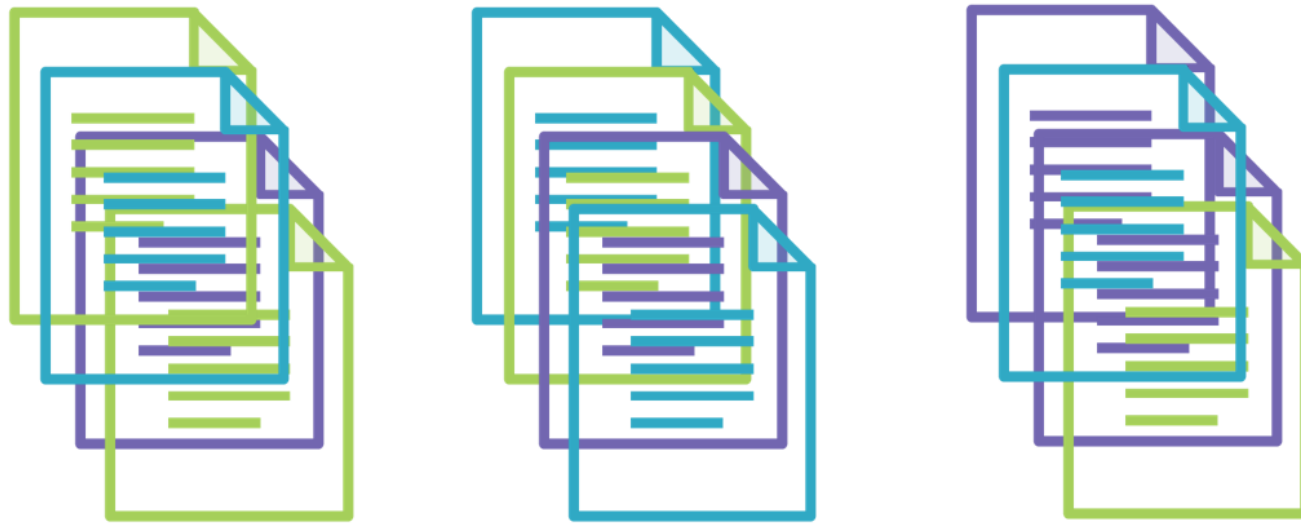
Has a unique name

Nodes join a cluster
using the cluster name



Basic Concepts of Elasticsearch

Document



**A whole bunch of documents that need to be
indexed so they can be **searched****

Basic Concepts of Elasticsearch

Document



catalog , reviews

- titles, description, comments

Basic Concepts of Elasticsearch

Type



Documents are divided into categories or **types**

Basic Concepts of Elasticsearch

Index



All of these types of documents make up an **index**

Basic Concepts of Elasticsearch

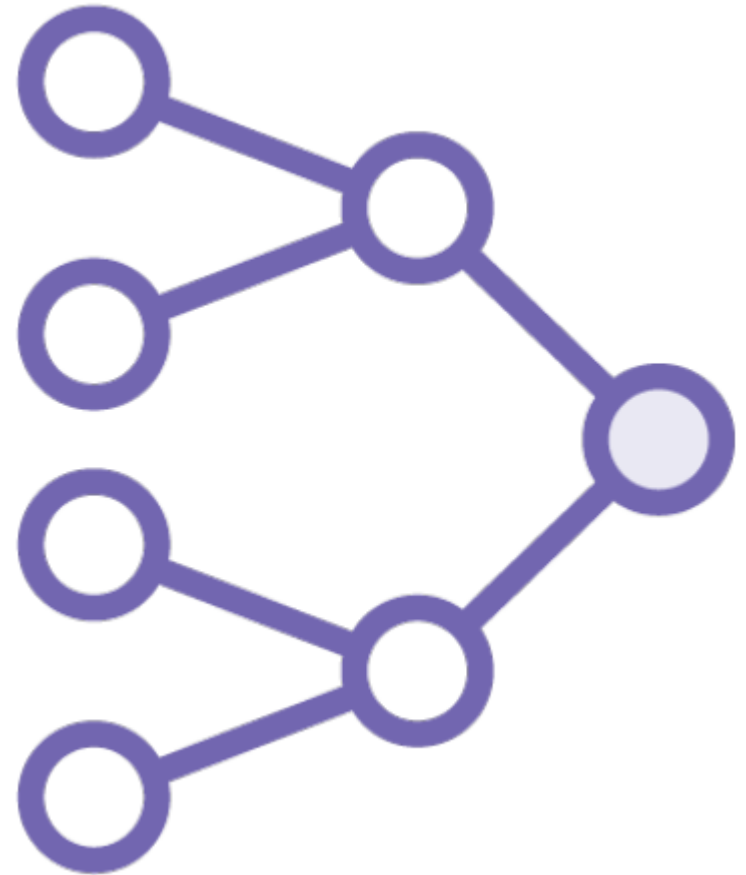
Index

Collection of similar documents

Identified by **name**

Any number of indices in a cluster

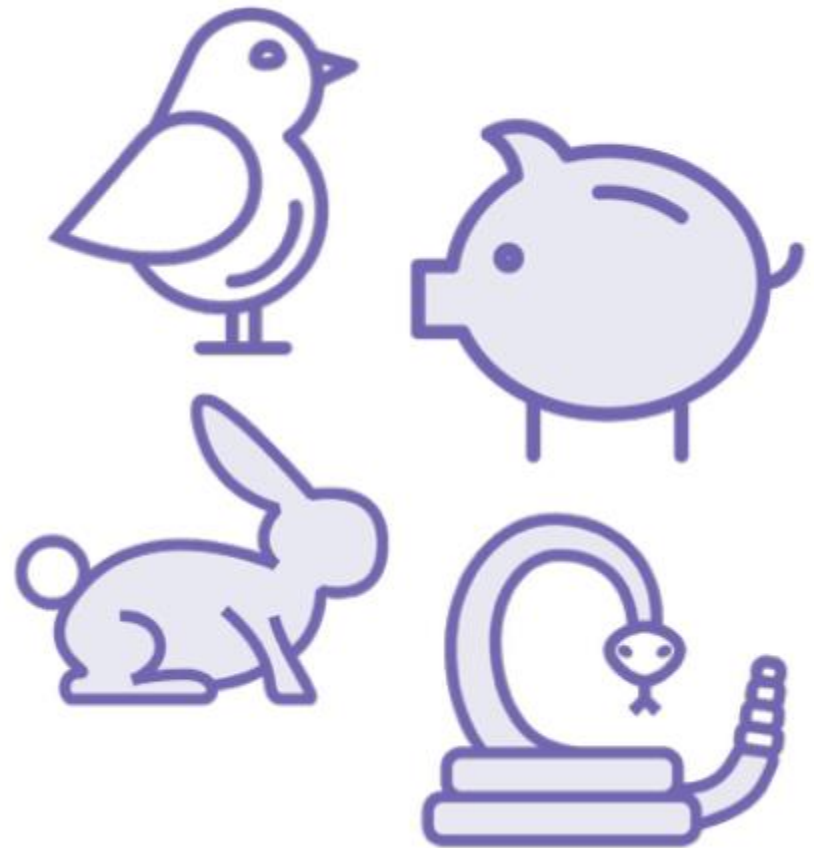
Different indices for different logical groupings



Basic Concepts of Elasticsearch

Type

Logical partitioning of documents
User defined grouping semantics
Documents with the same fields belong to one type



Basic Concepts of Elasticsearch

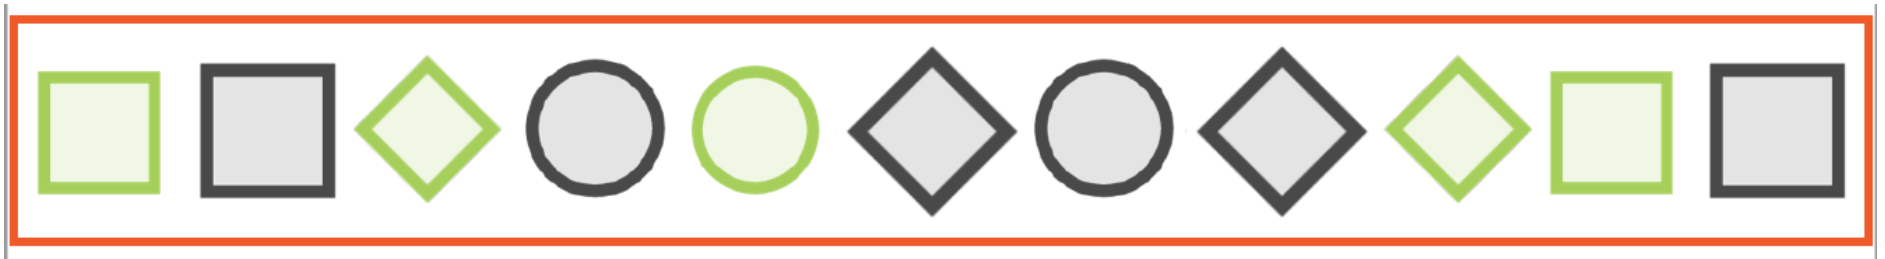
Document

- Basic unit of **information** to be indexed
- Expressed in **JSON**
- Resides within an index
- Assigned to a **type** within an index



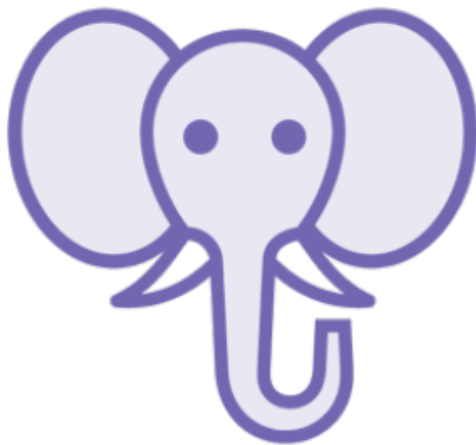
Index, Shards, Replicas

Document in an Index

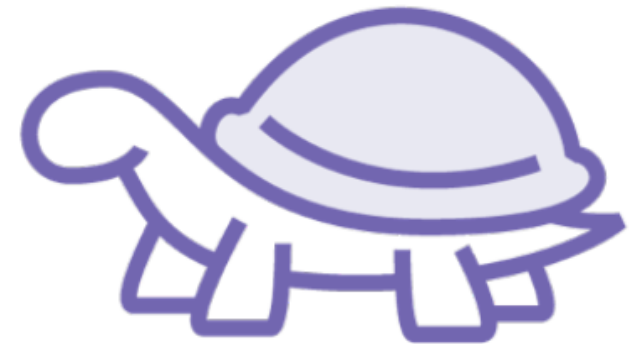


Index, Shards, Replicas

Document in an Index



Too **large** to fit in the
hard disk of one node



Too **slow** to serve all search
requests from one node

Index, Shards, Replicas

Shards



Split the index across multiple nodes in the cluster

Index, Shards, Replicas

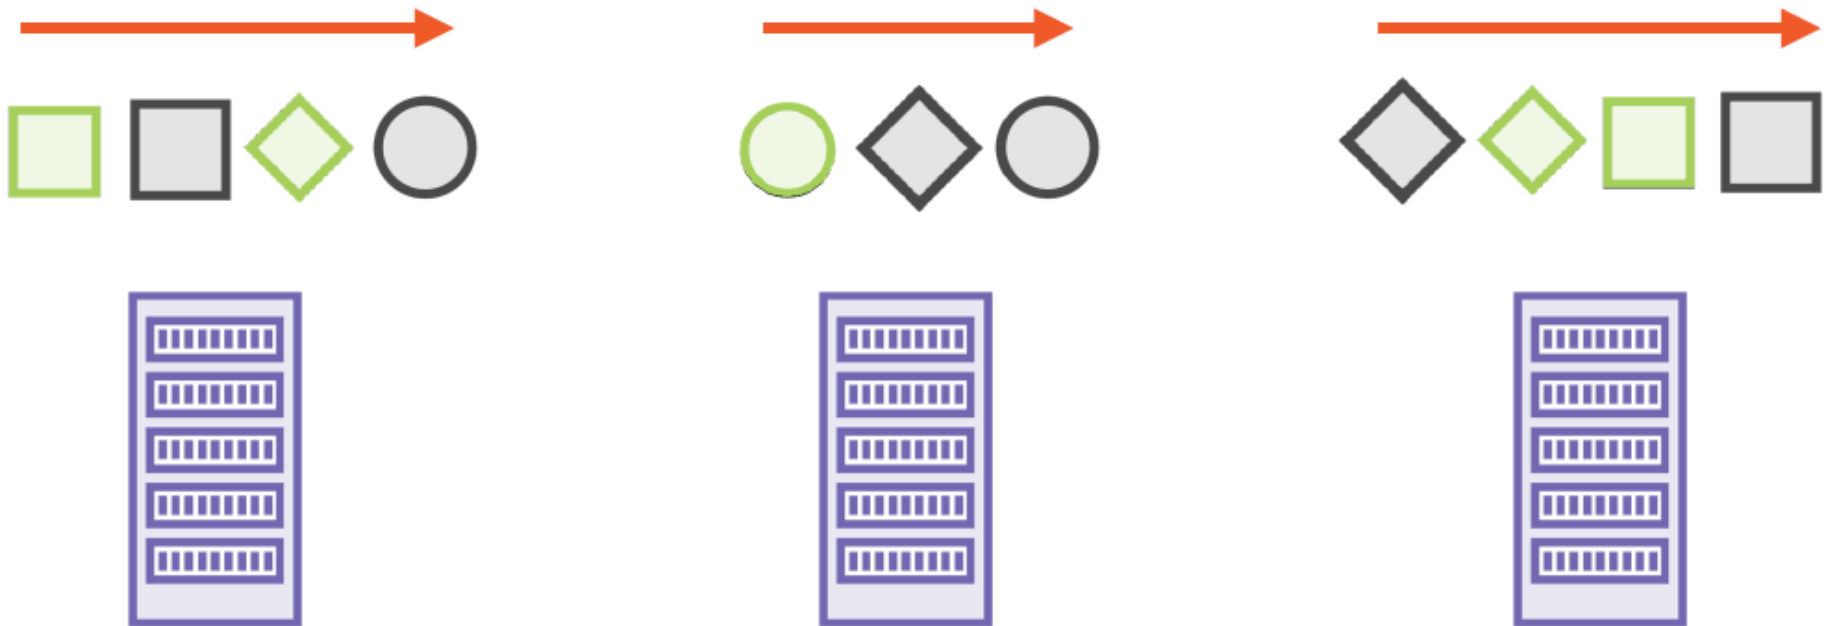
Shards



Sharding an index

Index, Shards, Replicas

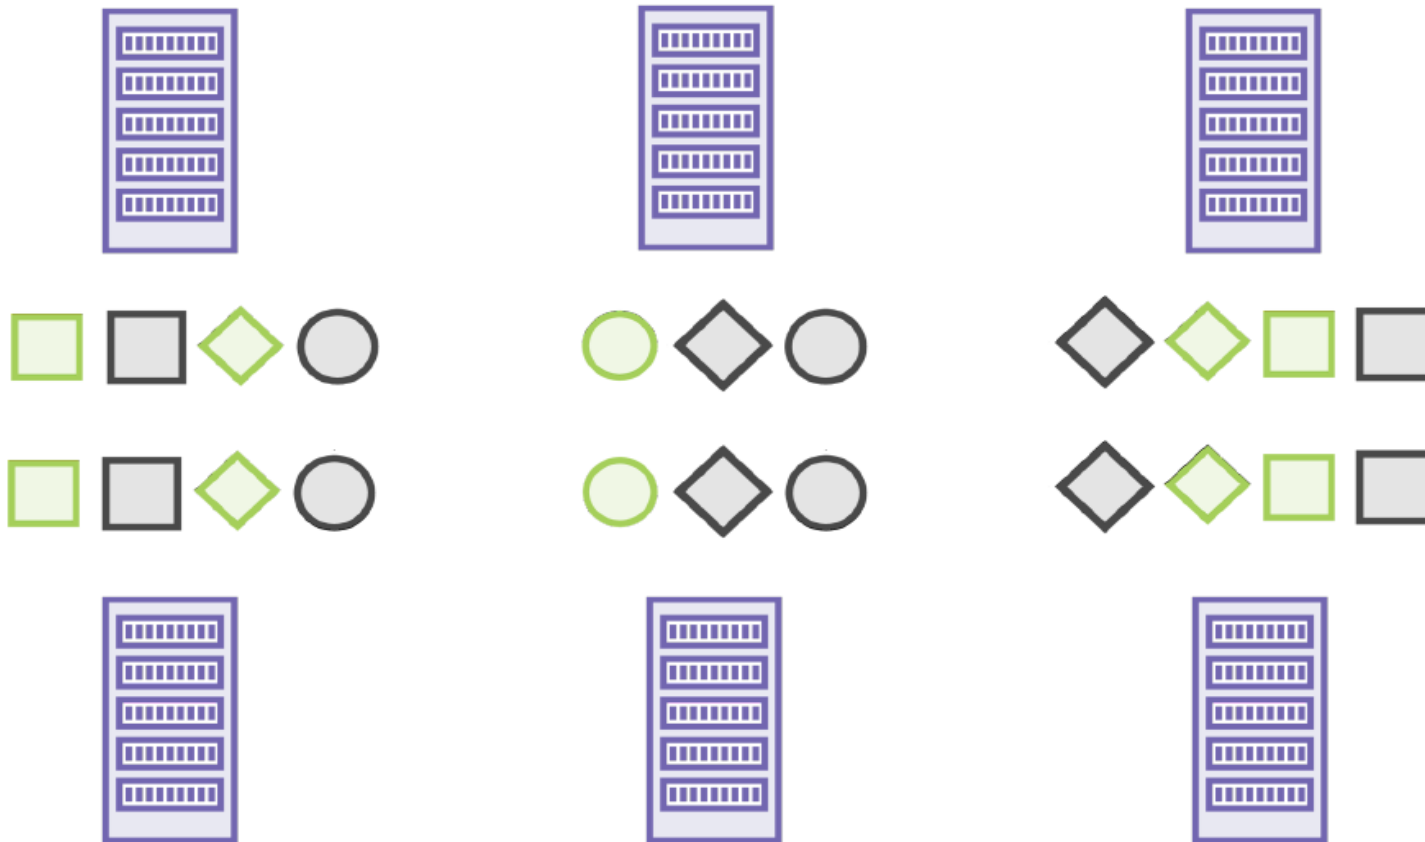
Shards



Search in parallel on multiple nodes

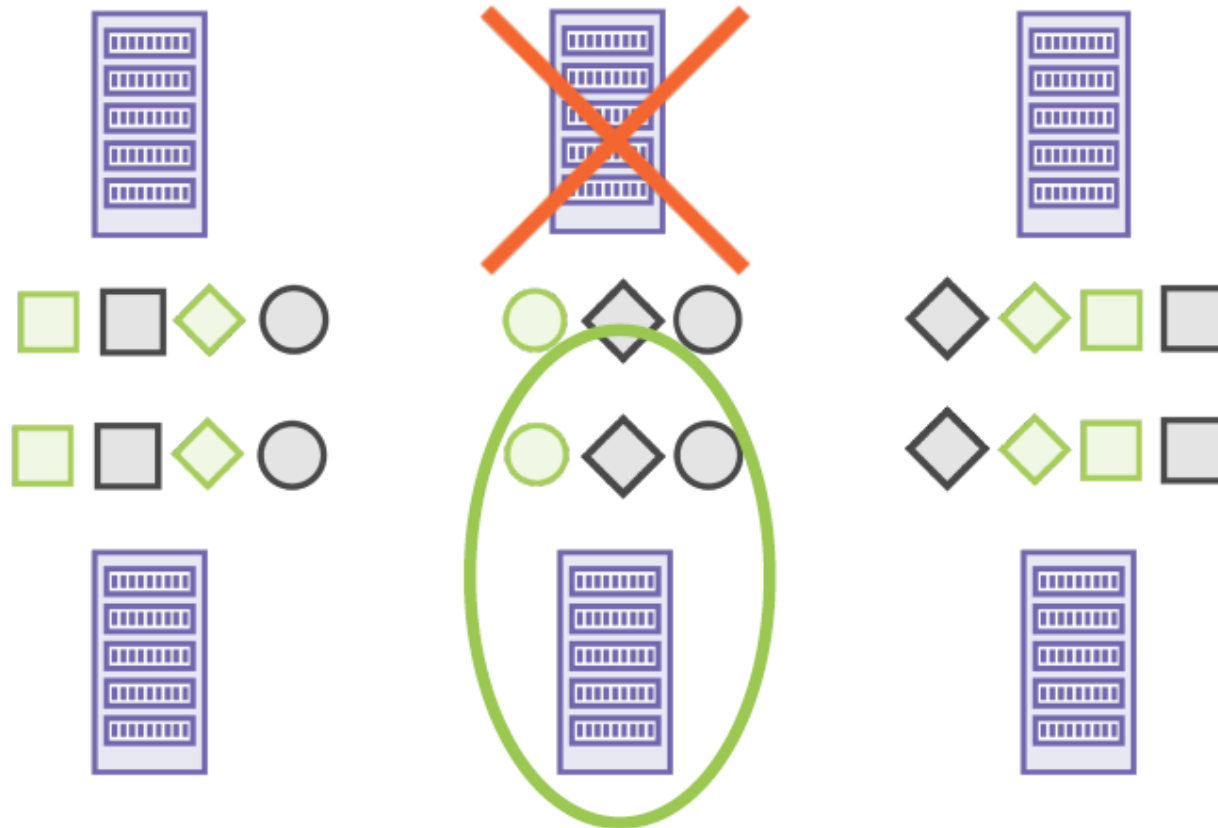
Index, Shards, Replicas

Replicas



Index, Shards, Replicas

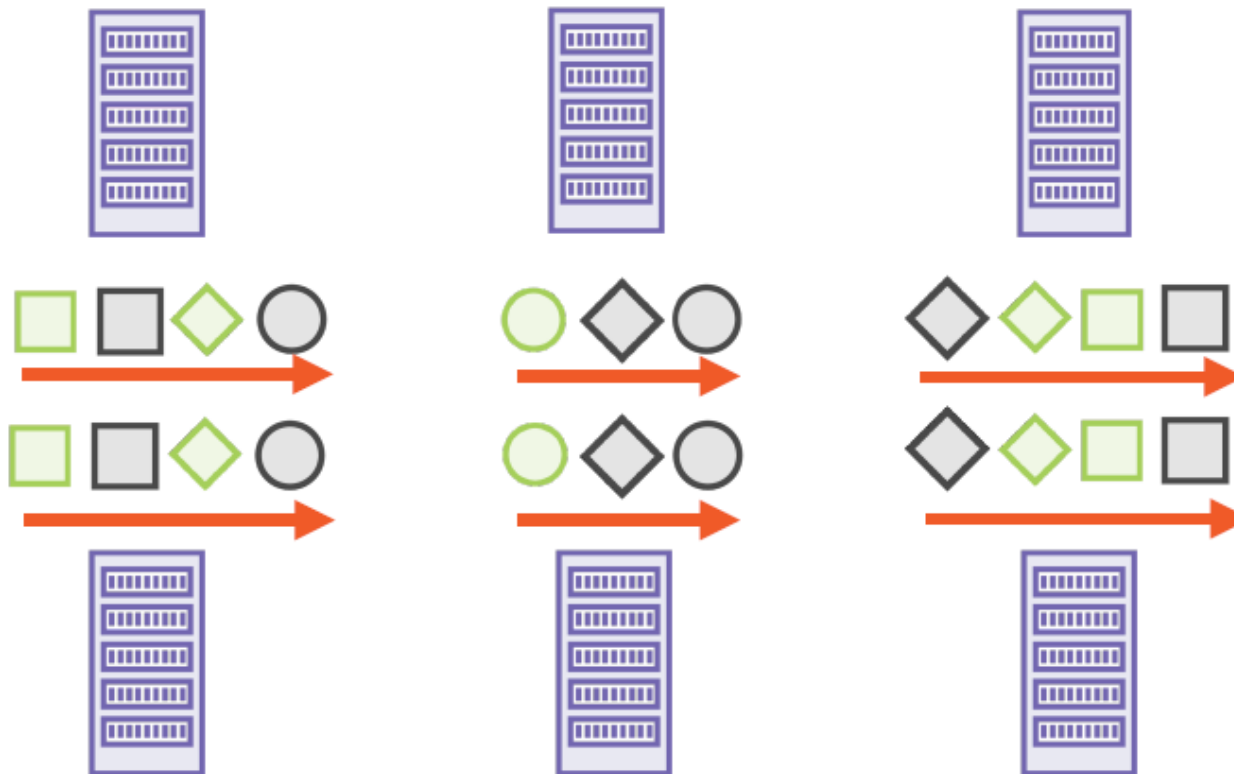
Replicas



High availability in case a node fails

Index, Shards, Replicas

Replicas

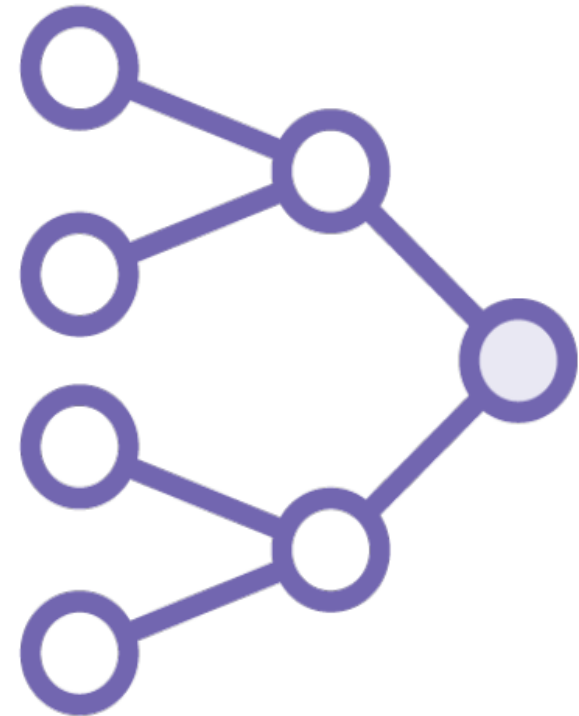


Scale search volume/throughput by searching multiple replicas

Index, Shards, Replicas

Shards and Replicas

- An index can be split into **multiple** shards
- A shard can be replicated **zero or more** times
- An index in Elasticsearch has 5 shards and 1 replica by default



Summary

- Learnt a little search engine history, ubiquitous nature of search
- Understood the basics steps involved in indexing and searching documents
- Learnt how the inverted index data structure works
- Got a brief introduction to Elasticsearch and its building blocks
- Set up and installed Elasticsearch on your local machine

Course Map – CRUD operations using the Elasticsearch APIs

- 1 RESTful APIs with Elasticsearch
- 2 Health and Index
- 3 CRUD
- 4 Bulk Operation on indexed document
- 5 Bulk Creation of indices from JSON data

RESTful APIs with Elasticsearch

RESTful APIs

- Elasticsearch uses REST APIs to administer the cluster, perform CRUD operations, search etc.
- Data is sent to and received from the server in JSON form



RESTful APIs with Elasticsearch

Cluster Health Status

```
curl "localhost:9200/_cat/health?v&pretty"
```

- **Green:**
 - All shards and replicas are available for requests, cluster fully functional
- **Yellow:**
 - Some replicas may not be available, cluster is still functional.
- **Red:**
 - Some shards not available, cluster NOT fully functional

RESTful APIs with Elasticsearch

cURL for Requests to REST APIs

- cURL is a tool which allows you to transfer data from and to a server using a variety of protocols
- HTTP, FTP, GOPHER, IMAP, LDAP etc.



CRUD

Demo

- Update documents by id:
 - whole documents
 - partial documents
- Delete documents in an index
- Delete the entire index

Bulk Operation on indexed document

Demo

- Bulk operations on documents:
 - retrieve multiple documents
 - index multiple documents
 - multiple operations in one command

Bulk Creation of indices from JSON data

Demo

- Bulk index documents from a JSON file

Summary

- Performed CRUD operations on indexes holding documents
- Implemented bulk operations on indexed documents
- Created indices in bulk from JSON data in a file

People matter, results count.



About Capgemini

With almost 180,000 people in over 40 countries, Capgemini is one of the world's foremost providers of consulting, technology and outsourcing services. The Group reported 2014 global revenues of EUR 10.573 billion.

Together with its clients, Capgemini creates and delivers business and technology solutions that fit their needs and drive the results they want. A deeply multicultural organization, Capgemini has developed its own way of working, the Collaborative Business Experience™, and draws on Rightshore®, its worldwide delivery model.



www.capgemini.com

